

CS528

SCO and Tuning


A Sahu
Dept of CSE, IIT Guwahati

Outline

- Intro to Code Optimization
- Machine independent/dependent optimization
- Common sense of Optimization
 - Do less work, avoid expensive Ops, shrink working set
- Simple measure Large impact : simd, branch, comm sub expre
- C++ Optimization
- Scalar Profiling
 - Manual Instrumentation (get_wall_time, clock_t)
 - Function and line based profiling (gprof, gcov)
 - Memory Profiling (valgrind, callgraph)
 - Hardware Performance Counter (oprofile,likwid)

CSO: Loop Jamming


```
for (i=0; i<10000; i++) {  
    Dostuff(i); //Small Independent work  
}  
for (i=0; i<10000; i++) {  
    DoMorestuff(i); //Small Independent work  
}
```



```
for (i=0; i<10000; i++) {  
    Dostuff(i);  
    DoMorestuff(i);  
}
```


CSO: Function Looping

```
for (i=0; i<10000; i++) {  
    Func (t, i) ;  
}  
Fun (int w, d) { //do lots of stuff }
```



```
funn (t) ;  
void funn (w) {  
    for (i=0; i<10000; i++) { //do lots stuffs  
        }  
}
```

CSO: Example: Vector ADT



```
vec_ptr new_vec(int len)
```

- Create vector of specified length

```
int get_vec_element(vec_ptr v, int  
index, int *dest)
```

- Retrieve vector element, store at *dest
- Return 0 if out of bounds, 1 if successful

```
int *get_vec_start(vec_ptr v)
```

- Return pointer to start of vector data

Optimization Example

- Procedure
 - Compute sum of all elements of vector
 - Store result at destination location
 - What's the Big-O of this code?

```
void combine1 (vec_ptr v, int *dest) {  
    int i;  
    *dest = 0;  
    for (i=0; i<vec_length(v); i++) {  
        int val;  
        get_vec_element(v, i, &val);  
        *dest += val;  
    }  
}
```

Move `vec_length` Call Out of Loop

- Value does not change from one iteration to next
- Code motion, `vec_length` requires only constant time, but significant overhead

```
void combine2 (vec_ptr v, int *dest) {  
    int i;  
    int length = vec_length(v) ;  
    *dest = 0;  
    for (i = 0; i < length; i++) {  
        int val;  
        get_vec_element(v, i, &val);  
        *dest += val;  
    }  
}
```

Reduction in Strength

```
void combine2(vec_ptr v, int *dest) {  
    int length = vec_length(v);  
    *dest = 0;  
    for (i = 0; i < length; i++) {  
        int val;  
        get_vec_element(v, i, &val);  
        *dest += val;  
    }  
}
```


Reduction in Strength

```
void combine3(vec_ptr v, int *dest) {  
    int i;  
    int length = vec_length(v);  
    int *data = get_vec_start(v);  
    *dest = 0;  
    for (i = 0; i < length; i++) {  
        *dest += data[i];  
    }  
}
```

Eliminate Unneeded Memory Refs

```
void combine4(vec_ptr v, int *dest) {  
    int i;  
    int length = vec_length(v);  
    int *data = get_vec_start(v);  
    int sum = 0;  
    for (i = 0; i < length; i++)  
        sum += data[i];  
    *dest = sum;  
}
```

Code Motion Example #2

- Procedure to Convert String to Lowercase
 - Extracted from many beginners' C programs
 - (Note: only works for ASCII, not extended characters)

```
void toLower(char *s) {  
    int i;  
    for (i = 0; i < strlen(s); i++)  
        if (s[i] >= 'A' && s[i] <= 'Z')  
            s[i] -= ('A' - 'a');  
}
```

Optimization Blocker: Procedure Calls

- *Why couldn't the compiler move `vec_len` or `strlen` out of the inner loop?*
 - Procedure might have side effects
 - Alters global state each time called
 - Function might not return same value for given arguments
 - Depends on other parts of global state
 - Procedure `lower` could interact with `strlen`

Optimization Blocker: Procedure Calls

- *Why doesn't compiler look at code for `vec_len` or `strlen`?*
 - Linker may overload with different version
 - Unless declared static
 - Interprocedural optimization is not extensively used, due to cost
- **Warning:**
 - Compiler treats procedure call as a black box
 - Weak optimizations in and around them

Profiling for Serial Code

Profiling for Serial Code

- Manual Instrumentation (`get_wall_time`, `clock_t`)
- Function and line based profiling (`gprof`, `gcov`)
- Memory Profiling (`valgrind`, `callgraph`)
- Hardware Performance Counter (`oprofile`, `likwid`)

Manual Instrumentation

- System Status
 - \$uptime, \$top , \$vmstat
 - \$systemmonitor, \$gnome-system-monitor
- \$time ./a.out
 - real time/wall clock time
 - cpu time and system time
 - cputime=sys time+usr time
- Using get_wall_time, clock_t

Manual Instrumentation

- \$time command and Using get_wall_time,

```
#include <time.h>

int main() {
    clock_t t; double Etime;
    t = clock();
    //Do some Work
    t = clock() - t;
    Etime= ((double) t) /CLOCKS_PER_SEC;
    printf("ETime =%f seconds", Etime)
    return 0;
}
```

Profiler: Hotspot Analyzer

- Given a program
- Finding out part of the program which takes maximum amount of time
- Optimizing hot-spot area reduce the execution time significantly
- Suppose a program spend 99% of time in a small function/code
 - Optimizing that code will result better performance

Function and line based profiling

- GNU profile (gprof)
 - `$gcc -p test.c`
 - `$/a.out`
 - `$gprof ./a.out`
 - `$gprof ./a.out >FPprofile.txt`
- GNU coverage (gcov)

Gprof Example

```
#include <stdio.h>

void FunA() {
    int i=0, g=0;
    while (i++<100000)
        { g+=i; }
}

void FunB() {
    int i=0, g=0;
    while (i++<400000)
        { g+=i; }
}
```

```
int main() {
    int iter=5000;
    while (iter--) {
        FunA();
        FunB();
    }
    return 0;
}
```

Gprof Example: Flat Profile

Flat profile:

Each sample counts as 0.01 seconds.

% cumulative	self	self	total	
time	seconds	seconds	calls	ms/call ms/call name
80.26	5.55	5.55	5000	1.11 1.11 FunB
20.94	6.99	1.45	5000	0.29 0.29 FunA

Gprof Example: Call Graph

Call graph

index	% time	self	children	called	name
				<spontaneous>	
[1]	100.0	0.00	6.99		main [1]
	5.55	0.00	5000/5000		FunB [2]
	1.45	0.00	5000/5000		FunA [3]

	5.55	0.00	5000/5000		main [1]
[2]	79.3	5.55	0.00	5000	FunB [2]

	1.45	0.00	5000/5000		main [1]
[3]	20.7	1.45	0.00	5000	FunA [3]

Function and line based profiling

- GNU profile (gprof)
- GNU coverage (gcov)
 - **`$gcc -fprofile-arcs -ftest-coverage tmp.c`**
 - **`$/a.out`**
 - **`$gcov tmp.c`**

File 'tmp.c'

Lines executed:87.50% of 8

Creating 'tmp.c.gcov'

Gcov output

```
#include <stdio.h>
int main (){
    int i, total;
    total = 0;
    for (i = 0; i < 10; i++)
        total += i;
    if (total != 45)
        printf ("Failure\n");
    else printf ("Success\n");
    return 0;
}
```

```
-: 1:#include <stdio.h>
1: 2:int main (){
-: 3:  int i, total;
1: 4:  total = 0;
11: 5:  for (i = 0; i < 10; i++)
10: 6:      total += i;
1: 7:  if (total != 45)
#####:8:      printf ("Failure\n");
1: 9:  else printf ("Success\n");
1: 10:  return 0;
-: 11:}
```


Valgrind

- Free tools: **\$sudo apt-get install valgrind**
- CallGraph, Profiler, Memory Check...
 - Many more
 - From C code, one can use API of valgrind
- Program analysis tools are useful
 - Bug detectors, Profilers, Visualizers
- **Dynamic binary analysis (DBA) tools**
 - Analyse a program's machine code at run-time
 - Augment original code with **analysis code**

Valgrind

```
void Work1(int n) {
    int i=0, j=0, k=0;
    while(i++<n) {
        while(j++<n) {while(k++<n) ;}
    }
}

void Work2(int n) { int i=0; while(i++<n) ;}
void Maneger(int n1, int n2) {
    Work1(n1); Work2(n2);
}

void Projects1() { Maneger(1000000, 1000);}
void Projects2() { Maneger(100, 1000000);}

int main() {
    Projects1(); Projects2(); return 0;
}
```

Valgrind: How to use

- `$gcc -pg -o Valgrindtest Valgrindtest.c`
- `$valgrind --tool=callgrind ./Valgrindtest`
- `$ls`

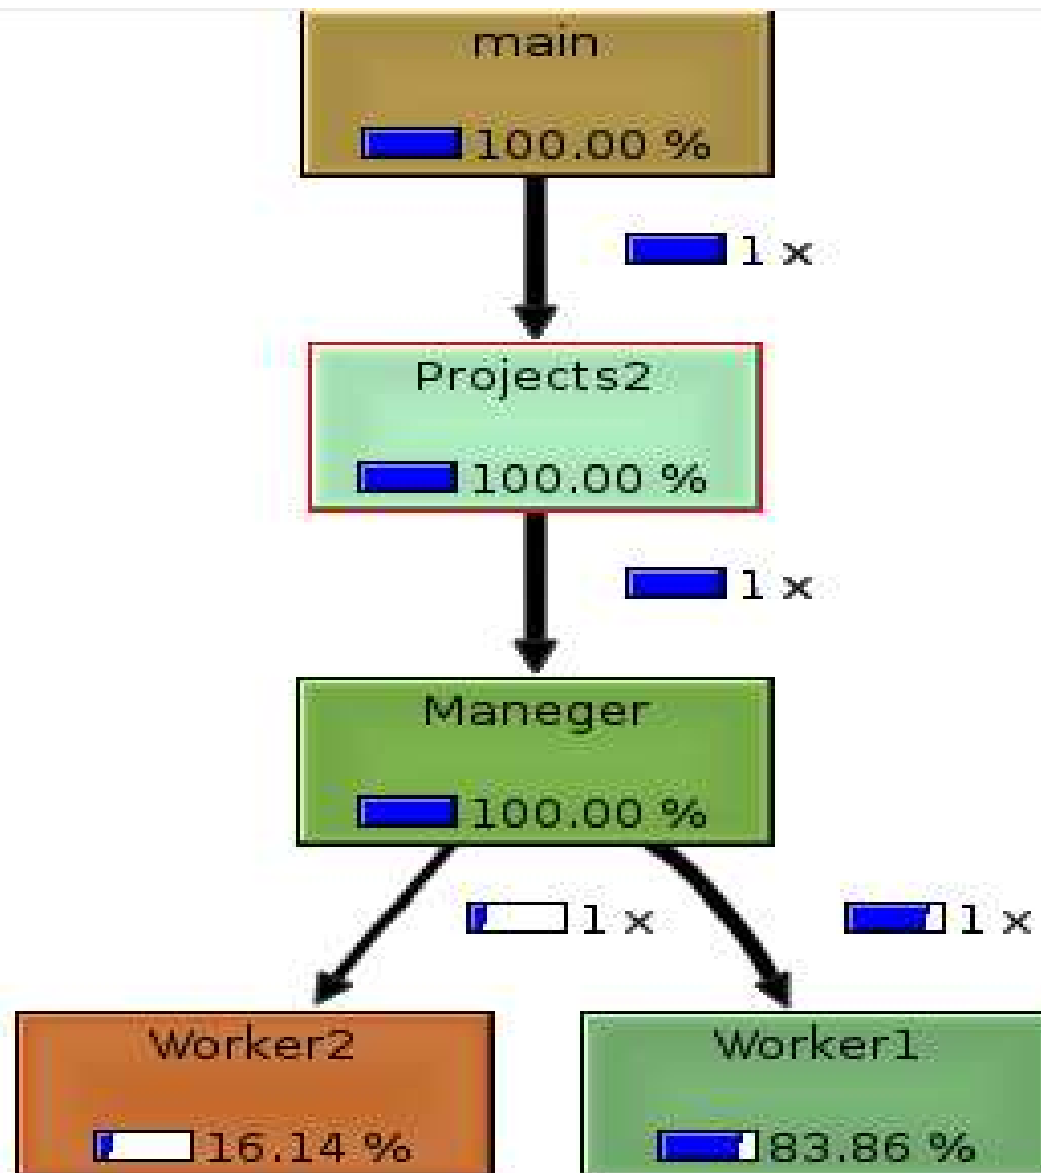
`Valgrindtest Valgrindtest.c callgrind.out.11233`

`$kcache-grind `ls -tr callgrind.out.* | tail -1``

pid



Valgrind: Call Graph



Further Optimizations for Serial Code

- Simple measure Large impact : simd, branch, comm sub expre
- C++ Optimization

Simple measures, large impact

- Elimination of Common Sub-expressions
- Avoid Branches:
 - Code Can be SIMDized by compiler/gcc
 - Effective use of pipeline for loop code
- Use of SIMD Instruction sets
 - 512 bit AVX SIMD in modern processor
 - ML/AI app use 8 bit Ops, can be speed up $512/8=64$ time by simply SIMD-AVX

Elimination of Common Sub-expressions

```
//value of s, r, x don't change in this loop
for (i=0; i<ALargeN; i++) {
    A[i]=A[i]+s+r+sinx(x);
}
```



```
//value of s, r, x don't change in this loop
Tmp=s+r+sinx(x);
for (i=0; i<ALargeN; i++) {
    A[i]=A[i]+Tmp;
}
```

Avoid Branches

```
for (i=0; i<N; i++)  
    for (j=0; j<N; j++) {  
        if (i<j) S=1; else S=-1;  
        C[i] =C[i]+S*A[i][j]*B[i];  
    }
```



```
for (i=0; i<N; i++) {  
    for (j=0; j<i; j++)  
        C[i] =C[i] -A[i][j]*B[i];  
    for (j=i; j<N; j++)  
        C[i] =C[i] +A[i][j]*B[i];  
}
```