

CS528

SIMD/GCC/FP

A Sahu
Dept of CSE, IIT Guwahati

Outline

- Intro to Code Optimization
- Machine independent/dependent optimization
- Common sense of Optimization
 - Do less work, avoid expensive Ops, shrink working set
- Simple measure Large impact : simd, branch, comm sub expre
- C++ Optimization
- Scalar Profiling
 - Manual Instrumentation (get_wall_time, clock_t)
 - Function and line based profiling (gprof, gcov)
 - Memory Profiling (valgrind, callgraph)
 - Hardware Performance Counter (oprofile,likwid)

Further Optimizations for Serial Code

- Simple measure Large impact : simd, branch, comm sub expre
- C++ Optimization

Simple measures, large impact

- Elimination of Common Sub-expressions
- Avoid Branches:
 - Code Can be SIMDized by compiler/gcc
 - Effective use of pipeline for loop code
- Use of SIMD Instruction sets
 - 512 bit AVX SIMD in modern processor
 - ML/AI app use 8 bit Ops, can be speed up $512/8=64$ time by simply SIMD-AVX

Elimination of Common Sub-expressions

```
//value of s, r, x don't change in this loop
for (i=0; i<ALargeN; i++) {
    A[i]=A[i]+s+r+sinx(x);
}
```



```
//value of s, r, x don't change in this loop
Tmp=s+r+sinx(x);
for (i=0; i<ALargeN; i++) {
    A[i]=A[i]+Tmp;
}
```

Avoid Branches

```
for (i=0; i<N; i++)  
    for (j=0; j<N; j++) {  
        if (i<j) S=1; else S=-1;  
        C[i] =C[i]+S*A[i][j]*B[i];  
    }
```



```
for (i=0; i<N; i++) {  
    for (j=0; j<i; j++)  
        C[i] =C[i] -A[i][j]*B[i];  
    for (j=i; j<N; j++)  
        C[i] =C[i] +A[i][j]*B[i];  
}
```

Use of SIMD: independent loop iteration

```
for (i=0; i<ALargeN; i++) {  
    A[i]=A[i]+B[i]*D[i];  
}
```

All iterations in this loop are independent : gcc SIMD utilize very nicely

//ML application uses 8 bit OPS, 512 bit AVX SIMD $512/8=64$ OPS can be done in parallel.

The ith iteration access : A[i], B[i], D[i]

Use of SIMD: independent loop iteration

$$S = \sum_{i=0}^N w_i x_i$$

- Vector dot product : is the most ***common and frequent kernel*** in
 - Matrix multiplication,
 - Neuron calculation (neural network NN)
 - Conv NN, Deep NN, //ML domain
 - Digital Signal Processing, Image Sig Processing, etc
 - Media Applications : audio, video, JPG/MPG, DCT.

Use of SIMD: independent loop iteration

$$S = \sum_{i=0}^N w_i x_i$$

```
for (i=0; i<ALargeN; i++) {  
    S=S+A[i]+B[i];  
}
```

All iterations in this loop are independent : gcc SIMD utilize very nicely

//ML application uses 8 bit OPS, 512 bit AVX SIMD 512/8=64 OPS can be done in parallel.

The ith iteration access : A[i], B[i]

Use of SIMD: independent loop iteration

```
for (i=0; i<N; i++) {  
    A[i]=A[i]+B[i]; //S1  
    B[i+1]=C[i]+D[i]; //S2  
}
```

The ith iteration access : A[i], B[i], B[i+1],C[i] D[i]

Dependent loop iteration : i and i+1


Use of SIMD: independent loop iteration

```
A[0]=A[0]+B[0];  
for (i=0; i<N; i++) {  
    B[i+1]=C[i]+D[i]; //S2  
    A[i+1]=A[i+1]+B[i+1]; //S1  
}  
B[N]=C[N-1]+D[N-1];
```

**The ith iteration access :
A[i+1], B[i+1],C[i], D[i]**

Use of SIMD: independent loop iteration

```
for (i=0; i<N; i++) {  
    A[i]=A[i]+B[i]; //S1  
    B[i+1]=C[i]+D[i]; //S2  
}
```



```
A[0]=A[0]+B[0];  
for (i=0; i<N; i++) {  
    B[i+1]=C[i]+D[i]; //S2  
    A[i+1]=A[i+1]+B[i+1]; //S1  
}  
B[N]=C[N-1]+D[N-1];
```

Use of SIMD: independent loop iteration

- Affine access : index $a.x+b$ form

```
for (i=0; i<N; i++) {  
    X[a*i+b]=X[c*i+d];  
    //where a,b,c,d are integer  
}
```

- $\text{GCD}(c,a)$ divides $(d-b)$ for loop dependence
- Ref Book : Hennesy Paterson, Advanced Computer Architecture, 5th Edition Book,

GCD test Example

```
for (i=0; i<N; i++) {  
    X[2*i+3]=X[2*i]+5.0;  
    //X[a*i+b]=X[c*i+d];  
}
```

- $\text{GCD}(c,a)$ must divide $(d-b)$ for loop dependence
- Value of $a=2$, $b=3$, $c=2$, $d=0$;
- $\text{GCD}(a,c)=2$, $d-b=-3$
- 2 does not divide -3 \rightarrow No dependence Possible

Role of Compilers

- General Compiler Optimization Options
- Inlining
- Aliasing
- Computational Accuracy

General Compiler Optimization Options

- GCC optimization : -O0, -O1, -O2,-O3
- \$man gcc
- At -O0 level:
 - Compiler refrain from most of the opt.
 - It is correct choice for analyzing the code with debugger
- At high level
 - Mixed up source lines, eliminate redundant variable, rearrange arithmetic expressions
 - Debugger has a hard time to give user a consistent view on code and data

General Compiler Optimization Options

- Level 1
 - fauto-inc-dec, -fmove-loop-invariant, -fmerge-constants, -ftree-copy-prop, -finline-fun-called-once
- Level 2
 - -falign-functions, -falign-loops, level, -finlining-small-fun, -finling-indirect-fun, -freorder-fun, -fstrict-aliasing
- Level 3
 - -ftree-slp-vectorize, -fvect-cost-model

Inlining

- Inlining
 - Tries to save overhead by inserting the complete code of function
 - At the place where it called
- Saves time and resources by
 - not using function call, stack
 - All compiler to use registers
 - Allows compiler to view a larger portion of code and employ OPTimization
- Auto inline or hint in program to function to be inlined

Aliasing

```
void scale_shift(double *a, double *b,  
                double s, int n) {  
    for(int i=1; i<n; i++)  
        a[i]=s*b[i-1];  
}
```

- Assuming a and b don't overlap
 - **double __restrict *a, double __restrict *b**
 - **__restrict say no overlap**
- Load and stores in the loop can be rearranged by compiler
- Apply software-pipelining, unrollling, group load/store, SIMD, etc

Computational Accuracy

- Compiler some time refrain from re-arranging arithmetic expression
- FP domain associative rule
 $a+(b+c) \neq (a+b)+c$
- If accuracy need to be maintained
 - Compared to non optimized code
 - Associative rules must not be used by compiler
 - Should be left to programmer to regroup safely
- FP underflow are push to zero

Computational Accuracy

- FP domain associative rule **$a+(b+c) \neq (a+b)+c$**
 - Let $a=1.0 \times 10^{38}$, $b=-1.0 \times 10^{38}$, $c=1$
 - Result of $a+(b+c)$
 - $= 1.0 \times 10^{38} + (-1.0 \times 10^{38} + 1)$
 - $= 1.0 \times 10^{38} + (-1.0 \times 10^{38}) = 0$ **//Big+Small=Big**
 - Result of $(a+b)+c$
 - $= (1.0 \times 10^{38} + -1.0 \times 10^{38}) + 1$
 - $= 0 + 1 = 1$

Computational Accuracy

- Why it happens for FP?
 - **FP format use 32 bit represent number up to $\pm 2^{127}$**
 - **Int use 32 bit represent up to $\pm 2^{31}$**
 - Used same 32 bit for large numbers, numbers are not equal-spaced
 - From 36000ft, both IITG and Amingoan are not distinguishable [Resolution:]
 - Going by Air: Delhi, Noida, Gurgaon use the same Airport