

CS528

Pthread and OpenMP

A Sahu

Dept of CSE, IIT Guwahati

Outline

- Threading
- OpenMP
- Cilk
- MPI

Multiprocessor Programming using Threading

Threading Language and Support

- Initially threading used for
 - Multiprocessing on single core : earlier days
 - Feeling/simulation of doing multiple work simultaneously even if on one processor
 - Used TDM, time slicing, Interleaving
- Now a day threading mostly used for
 - To take benefit of multicore
 - Performance and energy efficiency
 - We can use both TDM and SDM

Threading Language and Support

- Pthread: POSIX thread
 - Popular, Initial and Basic one
- Improved Constructs for threading
 - c++ thread : available in c++11, c++14
 - **Java thread : very good memory model**
 - Atomic function, Mutex
- Thread Pooling and higher level management
 - OpenMP (loop based)
 - Cilk (dynamic DAG based)

Pthread, C++ Thread, Cilk and OpenMP

```
pthread_t tid1, tid2;
pthread_create(&tid1, NULL, Fun1, NULL);
pthread_create(&tid2, NULL, Fun2, NULL);
pthread_join(tid1, NULL);
pthread_join(tid2, NULL);
```

Pthread

```
thread t1(Fun1);
thread t1(Fun2, 0, 1, 2); // 0, 1, 2 param to Fun2
t1.join();
t2.join();
```

C++
thread

```
#pragma omp parallel for
```

```
for(i=0;i<N;i++)
```

```
    A[i]=B[i]*C[i];
```

```
//Auto convert serial code to threaded code
```

```
// $gcc -fopenmp test.c; export OMP_NUM_THREADS=10; ./a.out
```

```
cilk fib (int n){//Cilk dynamic parallism, DAG recursive code
```

```
    if (n<2) return n;
```

```
    int x=spawn fib(n-1); //spawn new thread
```

```
    tnt y=spawn fib(n-2); //spawn new thread
```

```
    sync;
```

```
    return x+y;
```

```
}
```

Programming with Threads

- Threads
- Shared variables
- The need for synchronization
- Synchronizing with semaphores
- Thread safety and reentrancy
- Races and deadlocks

Traditional View of a Process

- Process = process context + code, data, and stack

Process context

Program context:

Data registers

Condition codes

Stack pointer (SP)

Program counter (PC)

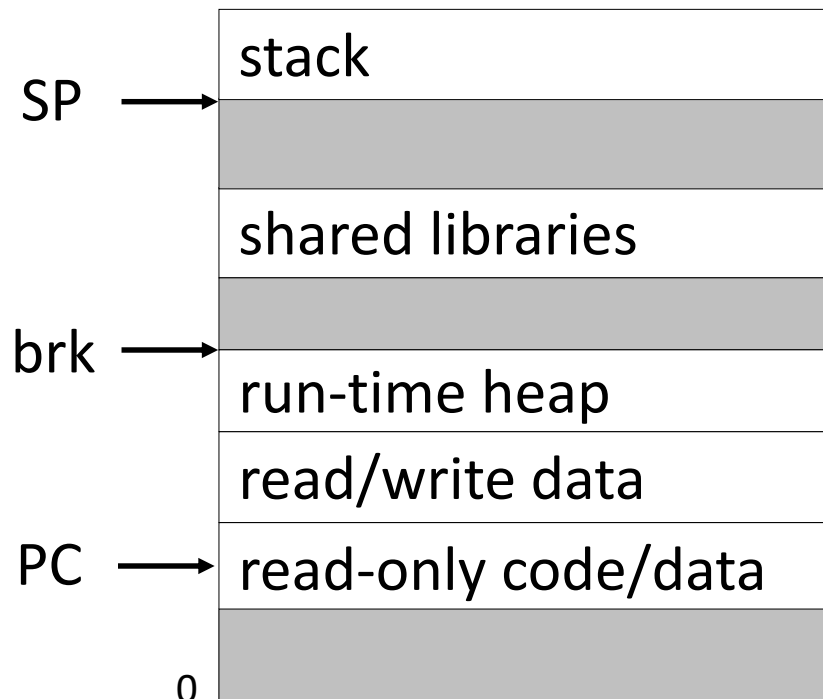
Kernel context:

VM structures (VMem)

Descriptor table

brk pointer

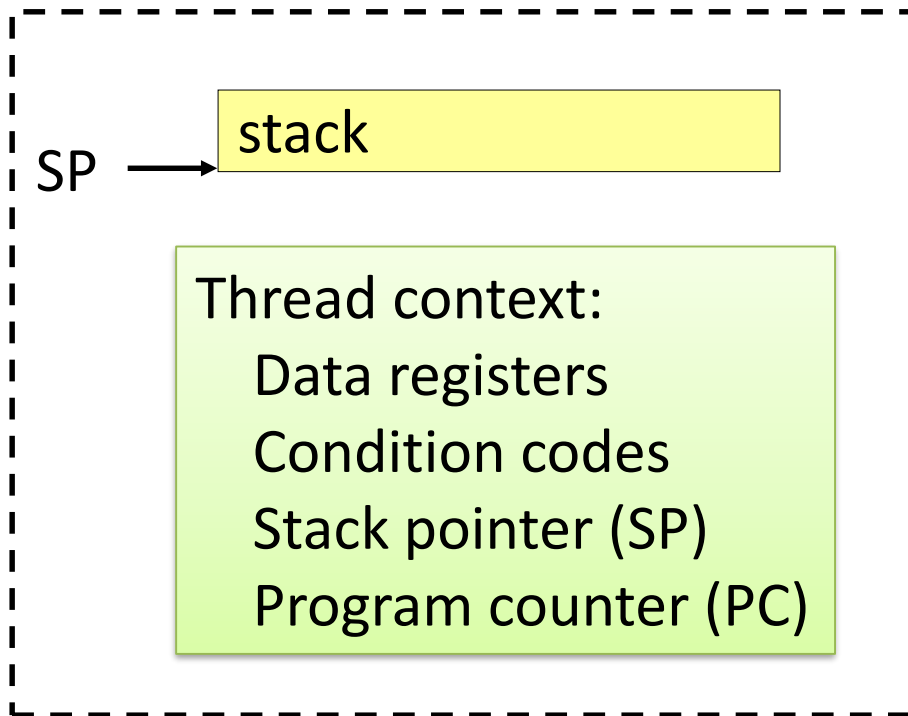
Code, data, and stack



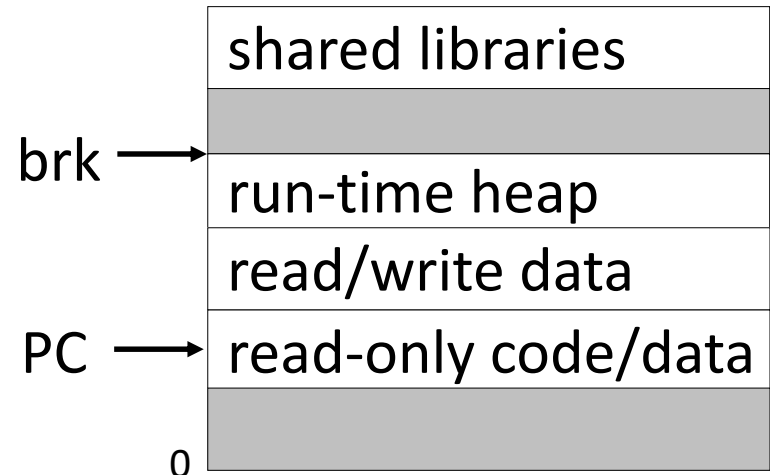
Alternate View of a Process

- Process = thread+ code, data & kernel context

Thread (main thread)



Code and Data



Kernel context:
VM structures
Descriptor table
brk pointer

A Process With Multiple Threads

- Multiple threads can be associated with a process
 - Each thread has its *own* logical control flow (sequence of PC values)
 - Each thread *shares* the same code, data, and kernel context
 - Each thread has its own thread id (TID)

A Process With Multiple Threads

Thread 1 (main thread)

stack 1

Thread 1 context:

Data registers

Condition codes

SP1

PC1

Shared code
and data

shared libraries

run-time heap

read/write data

read-only code/data

0

Kernel context:

VM structures

Descriptor table

brk pointer

Thread 2 (peer
thread)

stack 2

Thread 2 context:

Data registers

Condition codes

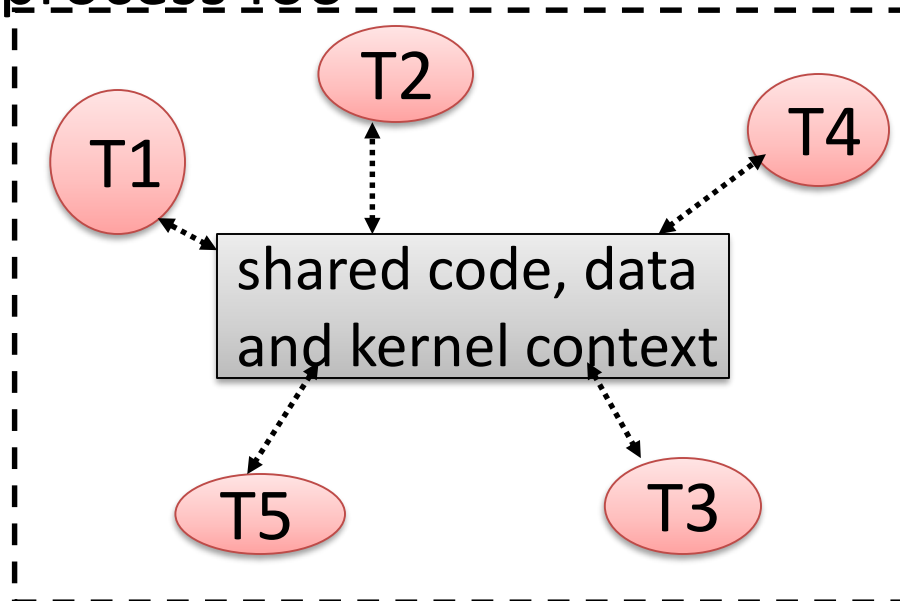
SP2

PC2


Logical View of Threads

- Threads associated with a process form a pool of peers
 - Unlike processes, which form a tree hierarchy

Threads associated with
process foo



Process hierarchy



Posix Threads (Pthreads) Interface

- **Creating and reaping threads**
 - `pthread_create`, `pthread_join`
- **Determining your thread ID** : `pthread_self`
- **Terminating threads**
 - `pthread_cancel`, `pthread_exit`
 - `exit` [terminates all threads], `return` [terminates current thread]
- **Synchronizing access to shared variables**
 - `pthread_mutex_init`,
`pthread_mutex_[un]lock`
 - `pthread_cond_init`,
`pthread_cond_[timed]wait`

The Pthreads "hello, world" Program

```
/* thread routine */  
void *HelloW(void *vargp) {  
    printf("Hello, world!\n");  
    return NULL;  
}
```

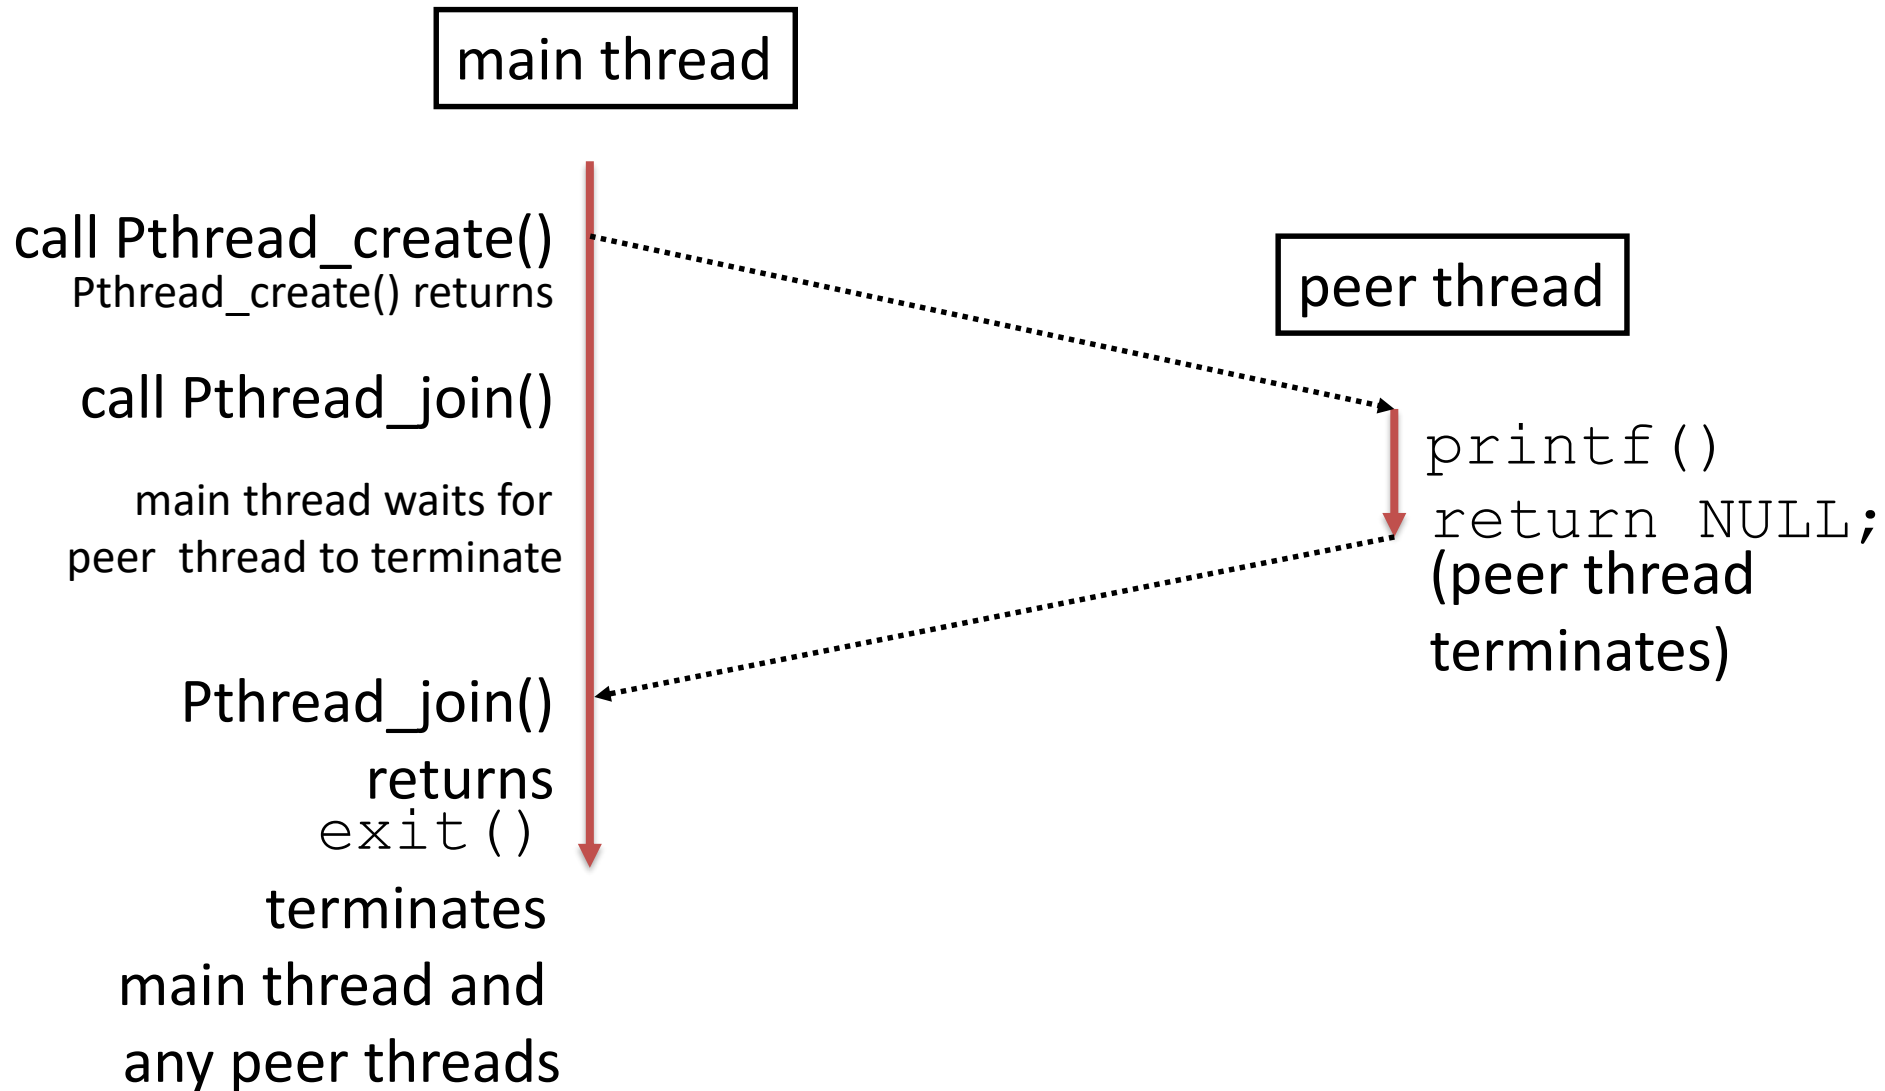
```
int main() {  
    pthread_t tid;  
    pthread_create(&tid, NULL, HelloW, NULL);  
    pthread_join(tid, NULL);  
    return 0;  
}
```

*Thread attributes
(usually NULL)*

*Thread arguments
(void *p)*

*return value
(void **p)*

Execution of Threaded “hello, world”



Pros and Cons: Thread-Based Designs

- **+ Easy to share data structures between threads**
 - E.g., logging information, file cache
- **+ Threads are more efficient than processes**
- **– Unintentional sharing can introduce subtle and hard-to-reproduce errors!**
 - Ease of data sharing is greatest strength of threads
 - Also greatest weakness!

VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];  
  
void VectorSumSerial() {  
    for( int j=0; j<SIZE; j++)  
        A[j]=B[j]+C[j];  
}
```

Suppose Size=1000

| | | | |
|--------------|----------------|----------------|----------------|
| 0-249 | 250-499 | 500-749 | 750-999 |
|--------------|----------------|----------------|----------------|

T1

T2

T3

T4

VectorSum Serial

```
int A[VSize], B[VSize], C[VSize];  
  
void VectorSumSerial() {  
    for( int j=0; j<SIZE; j++)  
        A[j]=B[j]+C[j];  
}
```

- Independent
- Divide work into equal for each thread
- Work per thread: $\text{Size}/\text{numThread}$

VectorSum Parallel

```
void *DoVectorSum(void *tid) {  
    int j, SzPerthrd, LB, UB, TID;  
    TID= *((int *)tid);  
    SzPerthrd=(VSize/NUM_THREADS);  
    LB= SzPerthrd*TID;UB=LB+SzPerthrd;  
  
    for (j=LB; j<UB; j++)  
        A[j]=B[j]+C[j];  
}
```

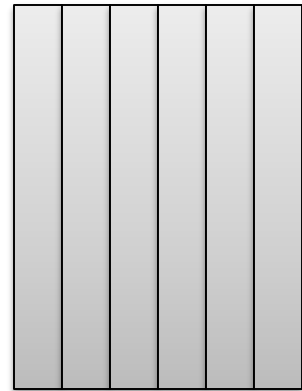
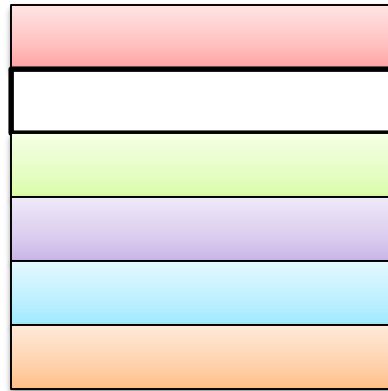
VectorSum Parallel

```
int main() {  
    int i;  
    pthread_t thread[NUM_THREADS];  
    for (i = 0; i < NUM_THREADS; i++)  
        pthread_create(&thread[i],  
            NULL, DoVectorSum, (void*)&i);  
    for (i = 0; i < NUM_THREADS; i++)  
        pthread_join(thread[i], NULL);  
  
    return 0;  
}
```

Matrix multiply and threaded matrix multiply

- Matrix multiply: $C = A \times B$

$$C[i, j] = \sum_{k=1}^N A[i, k] \times B[k, j]$$



Matrix multiply and threaded matrix multiply

- Matrix multiply: $C = A \times B$

$$C[i, j] = \sum_{k=1}^N A[i, k] \times B[k, j]$$

- Divide the whole rows to T chunks
 - Each chunk contains : N/T rows, Assume $N\%T=0$

Matrix multiply Serial

```
void MatMul () {  
    int i, j, k, S;  
    for (i=0; i<Size; i++)  
        for (j=0; j<Size; j++) {  
            S=0;  
            for (k=0; k<Size; k++)  
                S=S+A[i][k]*B[k][j];  
            C[i][j]=S;  
        }  
}
```

Matrix Pthreaded: RowWise

```
void * DoMatMulThread(void *arg) {  
    int i,j,k,S, LB,UB, TID, ThrdSz;  
    TID=*((int *)arg);ThrdSz=Size/NumThrd;  
    LB=TID*ThrdSz;UB=LB+ThrdSz;  
  
    for (i=LB;i<UB;i++)  
        for (j=0;j<Size;j++) {  
            S=0;  
            for (k=0;k<Size;k++)  
                S=S+A[i][k]*B[k][j];  
            C[i][j]=S;  
        }  
}
```


Matrix Pthreaded: RowWise

```
int main() {  
    pthread_t  thread[NumThread];  
    int t;  
    Initialize();  
    for (t=0; t<NumThread; t++)  
        pthread_create(&thread[t], NULL,  
            DoMatMulThread, &t);  
    for (t=0; t<NumThread; t++)  
        pthread_join(thread[t], NULL);  
    TestResult();  
    return 0;  
}
```

Estimating π using Monte Carlo


- The probability of a random point lying inside the unit circle:

$$\mathbf{P} \left(x^2 + y^2 < 1 \right) = \frac{A_{circle}}{A_{square}} = \frac{\pi}{4}$$

- If pick a random point N times and M of those times the point lies inside the unit circle:

$$\mathbf{P}^{\circ} \left(x^2 + y^2 < 1 \right) = \frac{M}{N}$$

- If N becomes very large, $\mathbf{P} = \mathbf{P}^{\circ}$



$$\pi = \frac{4 \cdot M}{N}$$

Value of PI: Monte-Carlo Method

```
void MontePI () {  
    int count=0,i;  
    double x,y,z;  
    for ( i=0; i<niter; i++) {  
        x = (double)rand()/RAND_MAX;  
        y = (double)rand()/RAND_MAX;  
        z = x*x+y*y;  
        if (z<=1) count++;  
    }  
    pi=(double) count/niter*4;  
}
```

PI- Multi-threaded

- 1 thread you are able to generate N points
 - Suppose M points fall under unit circle
 - $PI = 4M/N$
- With 10 thread generate 10XN points and calculate more accurately
 - Each thread calculate own value of PI (or M)
 - Average later on (or recalculate PI from collective M)

Value of PI: Pthreaded

```
int main() {  
    pthread_t  thread[NumThread]; double pi;  
    int t, at[NumThread], count, TotalIter;  
    for(t=0; t<NumThread; t++)  
        pthread_create(&thread[t], NULL,  
                        DoLocalMC_PI, &t);  
    for(t=0; t<NumThread; t++)  
        pthread_join(thread[t], NULL);  
    for(t=0; t<NumThread; t++) count+=LCount[t];  
    TotalIter=niter*NumThread;  
    pi=((double) count/TotalIter)*4;  
    return 0;  
}
```

Value of PI: Pthreaded

```
int LCount [NumThread];  
void *DoLocalMC_PI (void *aTid) {  
    int tid, count, i; double x, y, z;  
    tid= *((int *) aTid);  
    count=0; LCount[tid]=0;  
    for ( i=0; i<niter; i++) {  
        x = (double) rand() / RAND_MAX;  
        y = (double) rand() / RAND_MAX;  
        z = x*x+y*y; if (z<=1) count++;  
    }  
    LCount[tid]=count;  
}
```

Locking Shared Variable

Parallel Counter: without Lock

```
#define NITERS 100
int cnt = 0; /* shared */
int main() {
    pthread_t tid1, tid2;
    pthread_create(&tid1, NULL, count, NULL);
    pthread_create(&tid2, NULL, count, NULL);
    pthread_join(tid1, NULL);
    pthread_join(tid2, NULL);
    if (cnt != NITERS*2) printf("BOOM! cnt=%d", cnt);
    else printf("OK cnt=%d\n", cnt);
}
```

```
void *count(void *arg) {
for (int i=0; i<NITERS; i++) cnt++;
}
```

cnt should be 200
What went wrong?!

```
$ ./badcnt
BOOM! cnt=196
$ ./badcnt
BOOM! cnt=184
```


Thread Safety

- Functions called from a thread must be *thread-safe*
- There are four (non-disjoint) classes of thread-unsafe functions:
 - **Class 1: Failing to protect shared variables : L/UL**
 - Class 2: Relying on persistent state across invocations
 - Class 3: Returning pointer to static variable
 - Class 4: Calling thread-unsafe functions

Class 1: Failing to protect shared variables

- Fix: Use Lock and unlock semaphore operations
- Issue: Synchronization operations will slow down code
- Example: `goodcnt.c`

```
void *count(void *arg) {  
    for(int i=0;i<NITERS;i++)  
        pthread_mutex_lock(&LV);  
        cnt++;  
        pthread_mutex_unlock(&LV);  
} // LV is lock variable
```

Class 2: Relying on persistent state across multiple function invocations

- Random number generator relies on static state
- Fix: Rewrite function so that caller passes in all necessary state, → Maintain Thread Specific State

```
int rand() {  
    static uint next = 1;  
    next = next*1103515245 + 12345;  
    return (uint) (next/65536) % 32768;  
}  
  
void srand(uint seed) {  
    next = seed;  
}
```

Class 3: Returning pointer to static variable

- Fixes: 1. Rewrite code so caller passes pointer to `struct`, Issue: Requires changes in caller and callee
- *Lock-and-copy* : Issue: Requires only simple changes in caller (and none in callee), However, caller must free memory

```
struct hostent *gethostbyname (
                                char* name) {
    static struct hostent h;
    <contact DNS and fill in h>
    return &h;
}
```

Class 3: Returning pointer to static variable

```
struct hostent *gethostbyname_ts(char *p) {  
    struct hostent *q = Malloc(...);  
    P(&mutex); /* lock */  
    p = gethostbyname(name);  
    *q = *p;    /* copy */  
    V(&mutex);  
    return q;  
}
```

```
hostp = malloc(...);  
gethostbyname_r(name, hostp);
```

Class 4: Calling thread-unsafe functions

- Calling one thread-unsafe function makes an entire function thread-unsafe
- Fix: Modify the function so it calls only thread-safe functions

C++ Thread:atomic

```
atomic_uint AtomicCount;
```

```
void DoCount () {
```

```
    int j, timesperthrd;
```

```
    timesperthrd=(TIMES/NUM_THREADS);
```

```
    for (j=0; j<timesperthrd; j++) AtomicCount++;
```

```
}
```

```
main () {
```

```
    thread T[N_THRDS]; int i;
```

```
    for (i=0; i<N_THRDS; i++) T[i]=thread(DoCount);
```

```
    for (i=0; i<N_THRDS; i++) T[i].join();
```

```
}
```

Improved

```
atomic_uint AtomicCount;
```

```
void DoCount () {
```

```
    int j, timesperthrd, localcount=0;
```

```
    timesperthrd=(TIMES/NUM_THREADS);
```

```
    for (j=0; j<timesperthrd; j++) localcount++;
```

```
        AtomicCount+=localcount;
```

```
}
```

```
main () {
```

```
    thread T[N_THRDS]; int i;
```

```
    for (i=0; i<N_THRDS; i++) T[i]=thread(DoCount);
```

```
    for (i=0; i<N_THRDS; i++) T[i].join();
```

```
}
```


Java Thread:synchronized

```
void run()  
    synchronized(sender) {  
        // synchronizing the sender object //only one thread at time  
        sender.send(msg);  
    }  
}
```

```
synchronized void printTable(int n){//synchronized method  
    .....  
}
```

OpenMP

OpenMP

- Compiler directive: Automatic parallelization
- Auto generate thread and get synchronized

```
#include <openmp.h>
main() {
    #pragma omp parallel
    #pragma omp for schedule(static)
    {
        for (int i=0; i<N; i++) {
            a[i]=b[i]+c[i];
        }
    }
}
```

```
$ gcc -fopenmp test.c
```

```
$ export OMP_NUM_THREADS=4
```

```
$/a.out
```

OpenMP: Parallelism

Sequential code

```
for (int i=0; i<N; i++)  
    a[i]=b[i]+c[i];
```

OpenMP: Parallelism

(Semi) manual parallel

```
#pragma omp parallel
{
    int id =omp_get_thread_num();
    int Nthr=omp_get_num_threads();
    int istart = id*N/Nthr
    int iend= (id+1)*N/Nthr;
    for (int i=istart;i<iend;i++) {
        a[i]=b[i]+c[i];
    }
}
```

OpenMP: Parallelism

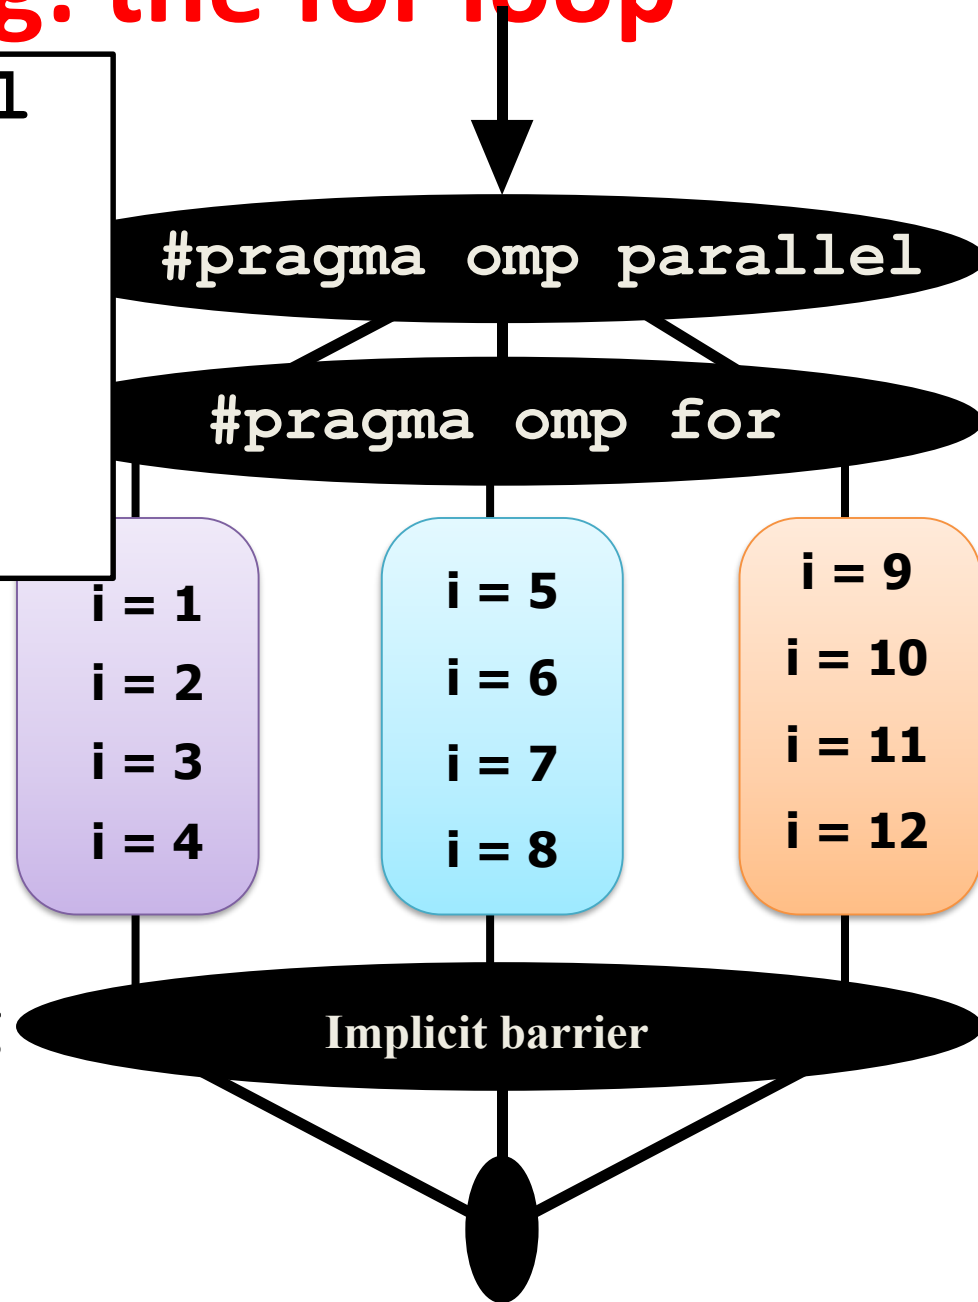
Auto parallel for loop

```
#pragma omp parallel
#pragma omp for schedule(static)
{
    for (int i=0; i<N; i++) {
        a[i]=b[i]+c[i];
    }
}
```

Work-sharing: the for loop

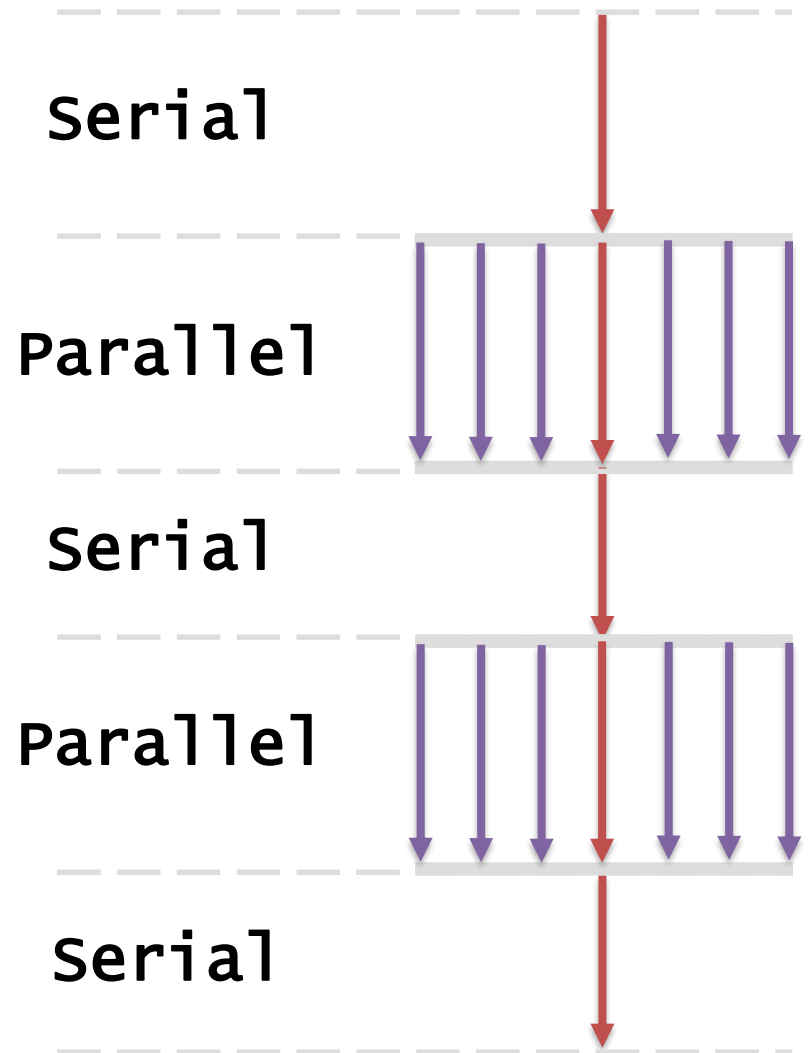
```
#pragma omp parallel
#pragma omp for
{
    for (i=1; i<13; i++)
        c[i]=a[i]+b[i];
}
```

- Threads are assigned an independent set of iterations
- Threads must wait at the end of work-sharing construct



OpenMP Fork-and-Join model

```
printf("begin\n");  
N = 1000;  
  
#pragma omp parallel for  
for (i=0; i<N; i++)  
    A[i] = B[i] + C[i];  
  
M = 500;  
  
#pragma omp parallel for  
for (j=0; j<M; j++)  
    p[j] = q[j] - r[j];  
  
printf("done\n");
```



AutoMutex: Critical Construct

```
sum = 0;
#pragma omp parallel private (lsum)
{
    lsum = 0;
    #pragma omp for
    for (i=0; i<N; i++) {
        lsum = lsum + A[i];
    }
    #pragma omp critical
    { sum += lsum; }
}
```


Threads wait their turn;
only one thread at a time
executes the critical section



Reduction Clause

Shared variable

```
sum = 0;  
#pragma omp parallel for reduction (+:sum)  
{  
    for (i=0; i<N; i++) {  
        sum = sum + A[i];  
    }  
}
```



OpenMP Schedule

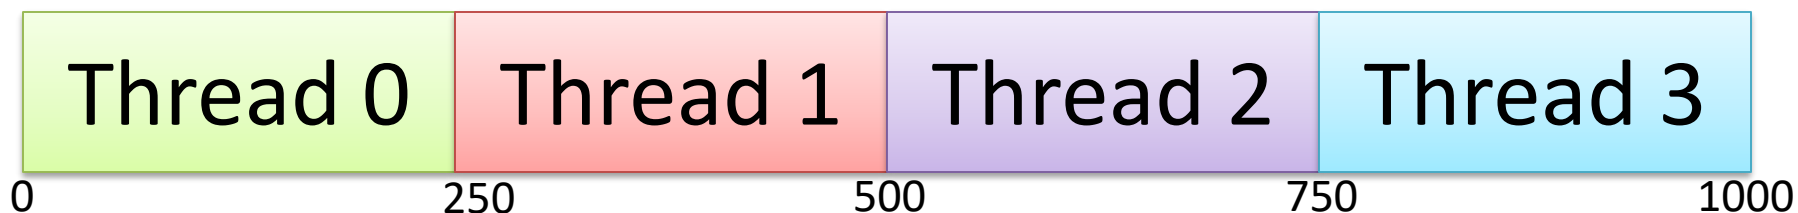
- Can help OpenMP decide how to handle parallelism

`schedule(type [,chunk])`

- **Schedule Types**
 - **Static** – Iterations divided into size chunk, if specified, and statically assigned to threads
 - **Dynamic** – Iterations divided into size chunk, if specified, and dynamically scheduled among threads

Static Schedule

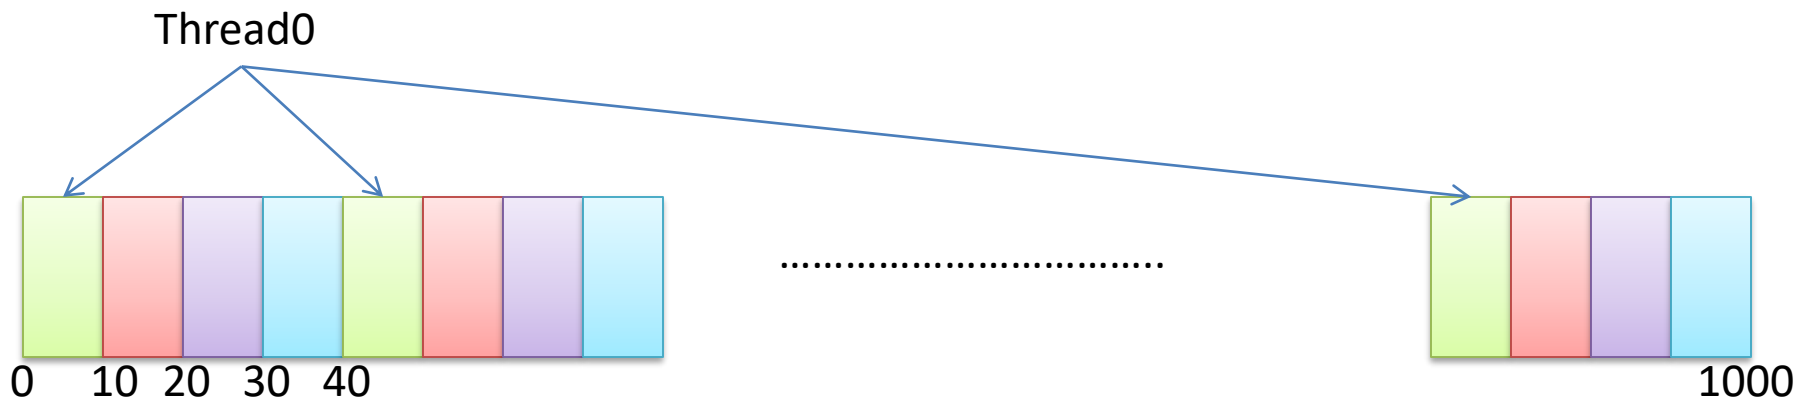
- Although the OpenMP standard does not specify how a loop should be partitioned
- Most compilers split the loop in N/p (N #iterations, p #threads) chunks by default.
- This is called a static schedule (with chunk size N/p)
 - *For example, suppose we have a loop with 1000 iterations and 4 omp threads. The loop is partitioned as follows:*



Static Schedule with chunk

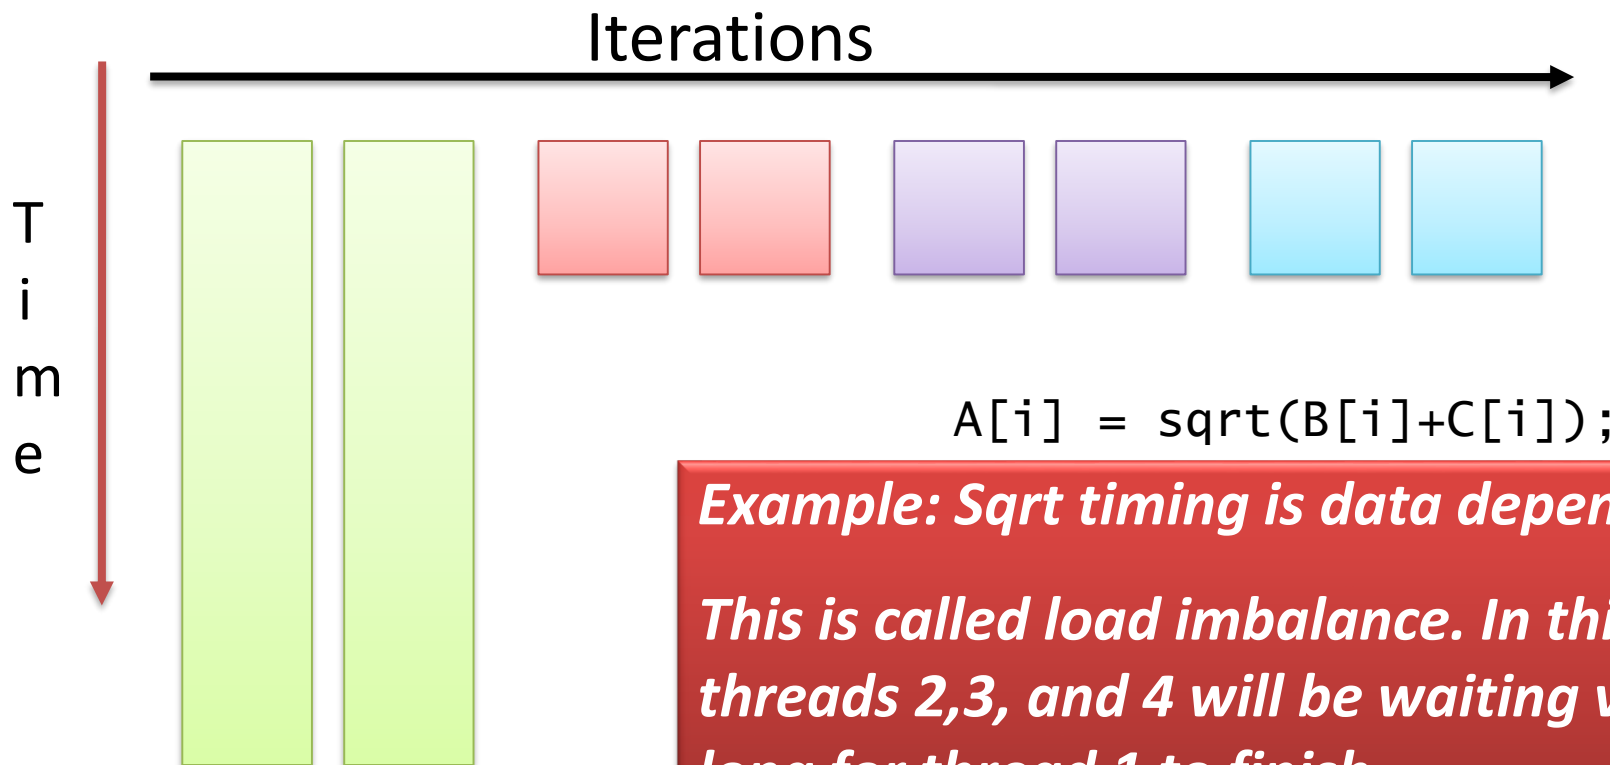
- A loop with 1000 iterations and 4 omp threads. Static Schedule with Chunk 10

```
#pragma omp parallel for schedule (static, 10)
{
  for (i=0; i<1000; i++)
    A[i] = B[i] + C[i];
}
```



Issues with Static schedule

- With static scheduling the number of iterations is evenly distributed among all openmp threads (i.e. Every thread will be assigned similar number of iterations).
- This is not always the best way to partition. Why is This?



Example: Sqrt timing is data dependent...

This is called load imbalance. In this case threads 2,3, and 4 will be waiting very long for thread 1 to finish

Dynamic Schedule

- With a dynamic schedule new chunks are assigned to threads when they come available.
- SCHEDULE(DYNAMIC,n)
 - Loop iterations are divided into pieces of size chunk. When a thread finishes one chunk, it is dynamically assigned another.

Dynamic Schedule

- SCHEDULE(GUIDED,n)
 - Similar to DYNAMIC but chunk size is relative to number of iterations left.
- Although Dynamic scheduling might be the preferred choice to prevent load imbalance
 - In some situations, there is a significant overhead involved compared to static scheduling.

More Examples on OpenMP

- <http://users.abo.fi/mats/PP2012/examples/OpenMP/>