

CS528

Scheduling of Dependent Tasks

A Sahu

Dept of CSE, IIT Guwahati

Outline

- $P_m \mid p_j, \text{pmtn} \mid C_{\max}$: Linear time solution
- $Q_m \mid p_j, \text{pmtn} \mid C_{\max}$: Poly time solution
- $Q_m \mid \text{ptmn} \mid \sum C_j$ Optimal Solution
- $P_m \mid p_j \mid C_{\max}$
 - **ILP Solution : Exponential**
 - 2 Approx, $2-1/m$ approx.
 - LPT : $3/2$ and $4/3$ Approx
- $P_m \mid p_j=1 \mid \sum w_j U_j$ Optimal Solution
- $P_m \mid p_j \mid \sum U_j$ NPC, Heuristic and Counter example
- $P_m \mid \text{pmtn}, p_j \mid \sum U_j$ **in NPC**
- $P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$ **in NPC**
 - 2 Approx

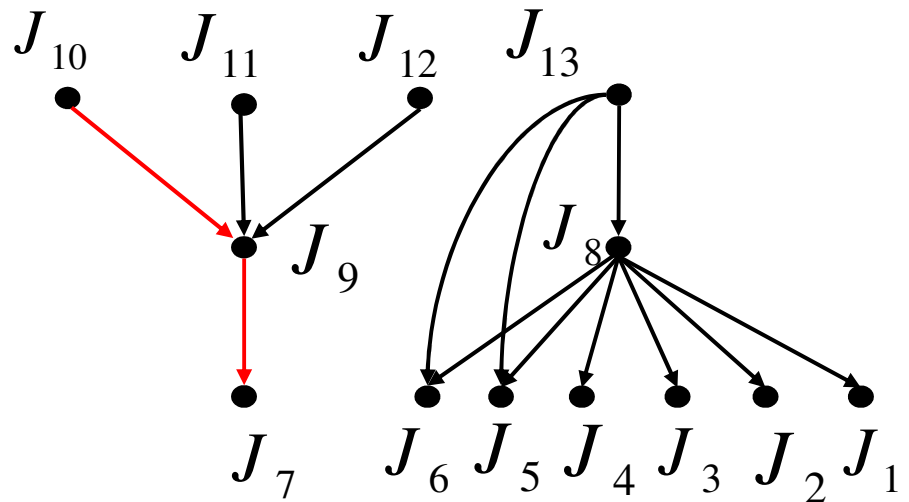
Scheduling of Dependent Tasks

Precedence constraints (*prec*)

Before certain jobs are allowed to start processing, one or more jobs first have to be completed.

Definition

- Successor
- Predecessor
- Immediate successor
- Immediate predecessor
- Transitive Reduction



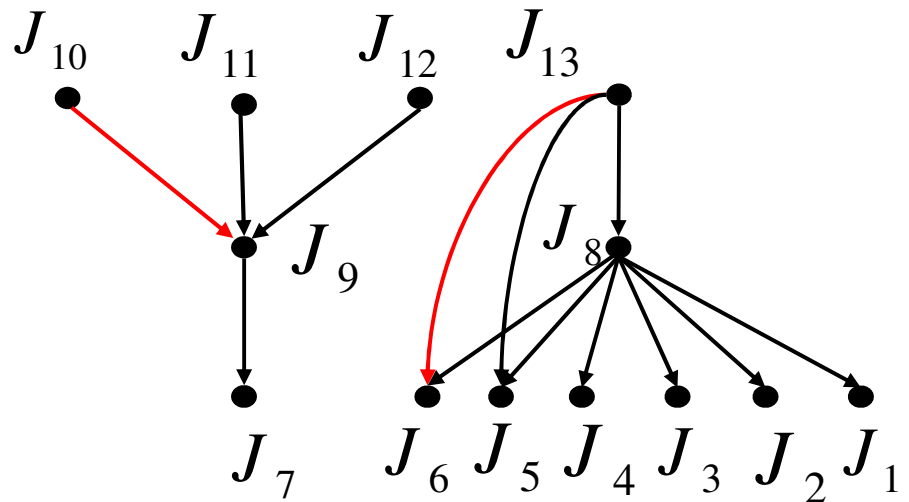
$$p(J_i) = 1$$

Precedence constraints (*prec*)

One or more job have to be completed before another job is allowed to start processing.

Definition

- Successor
- Predecessor
- Immediate successor
- Immediate predecessor
- Transitive Reduction



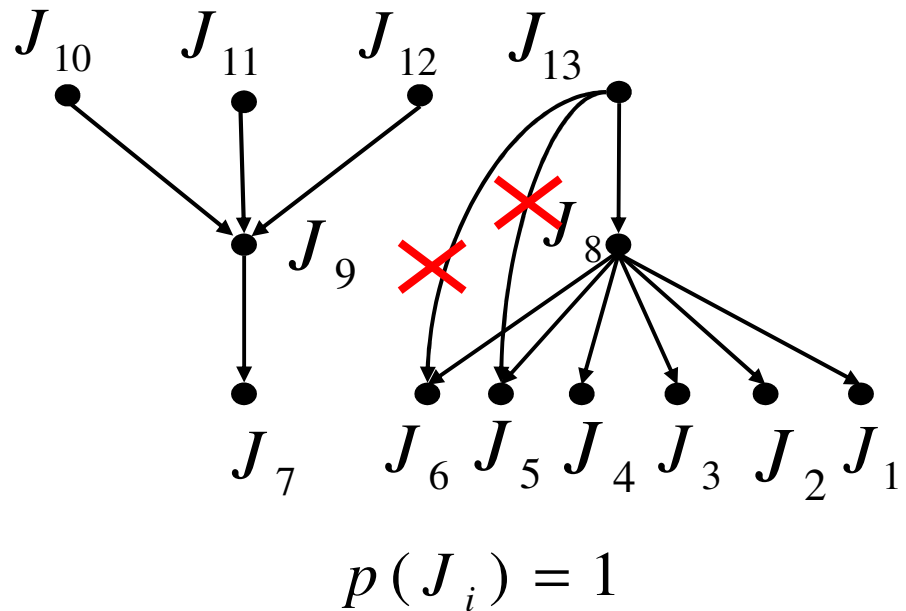
$$p(J_i) = 1$$

Precedence constraints (*prec*)

One or more job have to be completed before another job is allowed to start processing. *Prec : Arbitrary acyclic graph*

Definition

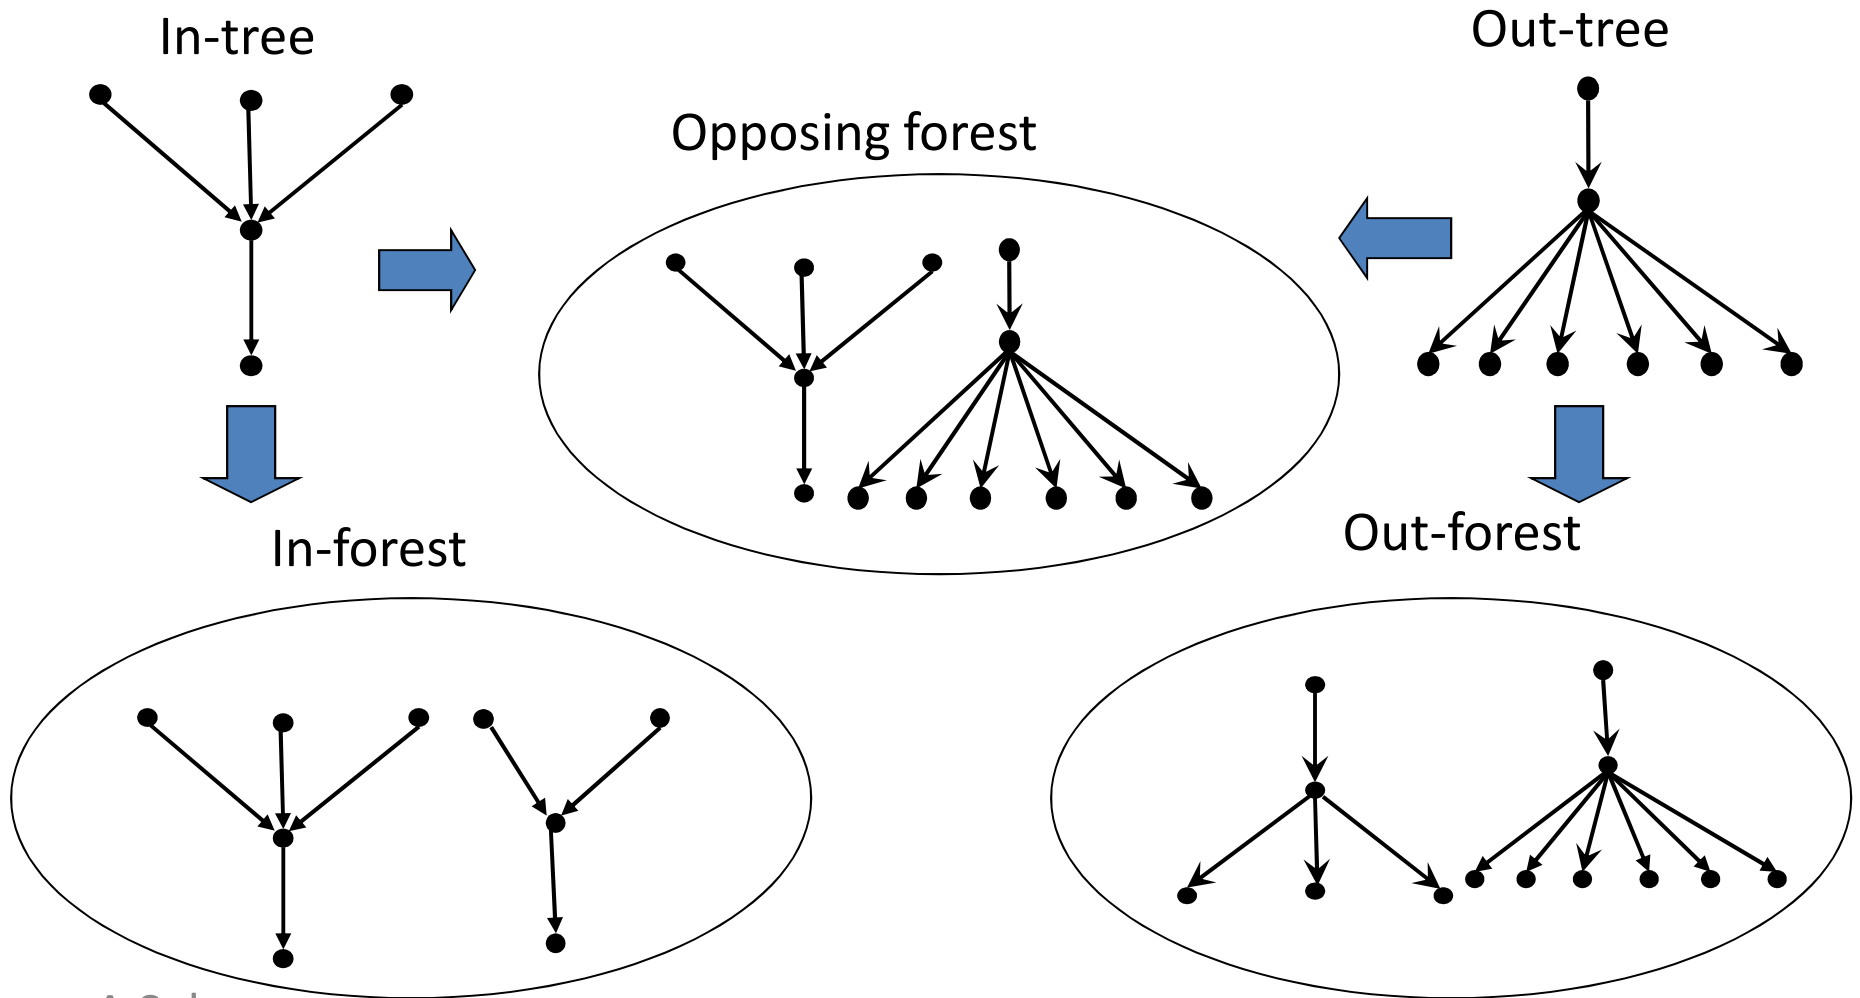
- Successor
- Predecessor
- Immediate successor
- Immediate predecessor
- **Transitive Reduction**



Special precedence constraints

- In-tree (Out-tree)
- In-forest (Out-forest)
- Opposing forest
- *Interval orders*
- *Series-parallel orders*
- *Level orders*

Special precedence constraints



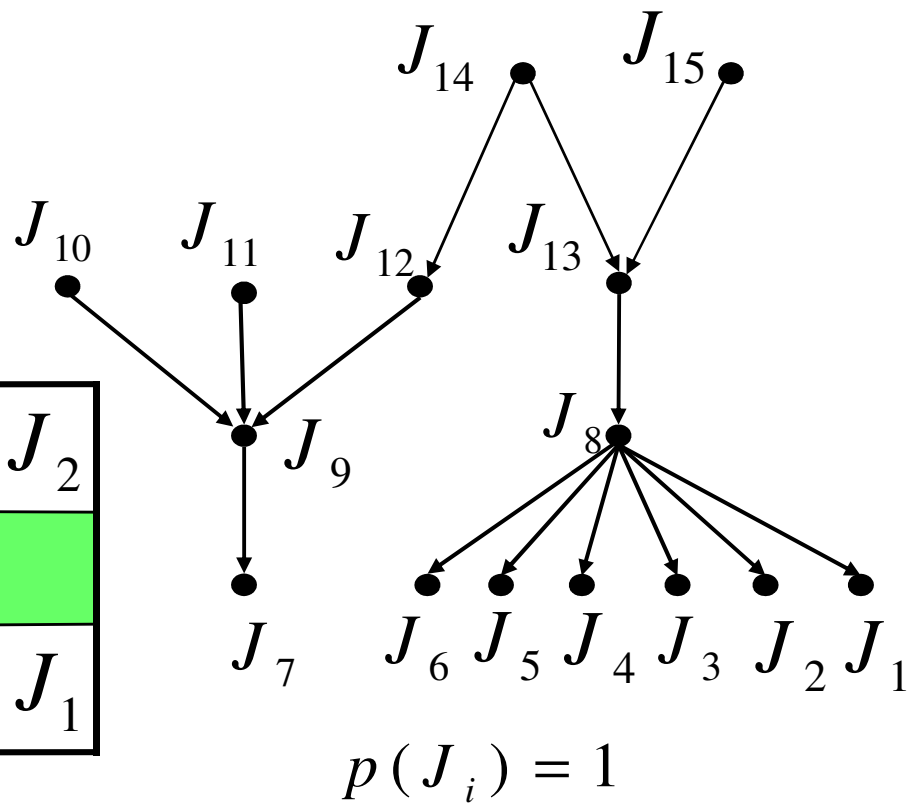
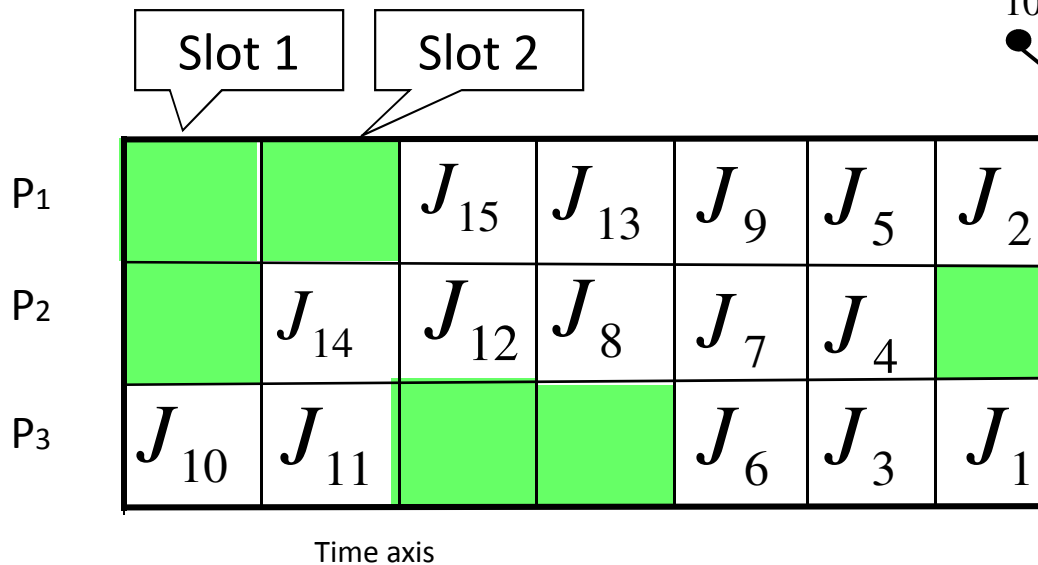
$P_m \mid \text{prec}, p_j = 1 \mid C_{\max} (m \geq 1)$

- Processor Environment
 - m identical processors are in the system.
- Job characteristics
 - Precedence constraints are given by a precedence graph;
 - Preemption is not allowed;
 - The release time of all the jobs is 0.
- Objective function
 - C_{\max} : the time the last job finishes execution.
 - If c_j denotes the finishing time of J_j in a schedule S ,

$$C_{\max} = \max_{1 \leq j \leq n} c_j$$

Gantt Chart

A Gantt chart indicates the time each job spends in execution, as well as the processor on which it executes *of some Schedule*



$$P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$$

Theorem 1

$P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$ is NP-complete.

1. Ullman (1976)

$$3\text{SAT} \leq P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$$

2. Lenstra and Rinnooy Kan (1978)

$$k\text{-clique} \leq P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$$

$P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$ is NP-complete.

Proof: out of Syllabus

$P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$

Mayr (1985)

- **Theorem 2**

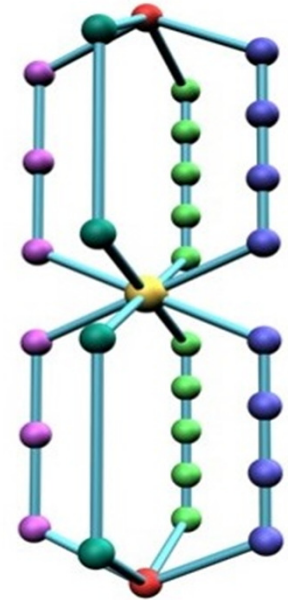
$P_m \mid p_j = 1, SP \mid C_{\max}$ is NP-complete.

SP: Series - parallel

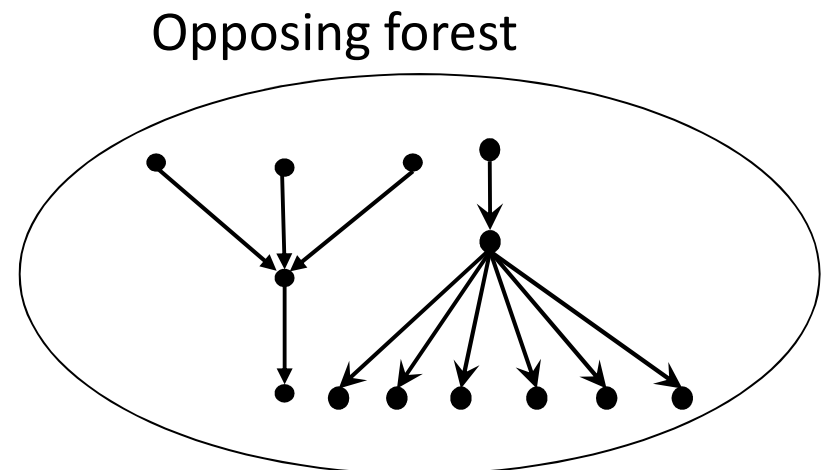
- **Theorem 3**

$P_m \mid p_j = 1, OF \mid C_{\max}$ is NP-complete.

OF: Opposing - forest




Proof: out of Syllabus



Hu's HLF/CP Algorithm

- T. C. Hu (1961), **Critical Path/Highest Level First**
- Assign a level h to each job.
 - If job has no successors, $h(j)$ equals 1.
 - Otherwise, $h(j)$ equals one plus the maximum level of its immediate successors.
- Set up a priority list L by nonincreasing order of the jobs' levels.
- Execute the list scheduling policy on this level based priority list L .

HLF/CP algorithm : Example



M2	J_{10}	J_{13}	J_8	J_6	J_3
M2	J_{11}	J_9	J_7	J_5	J_2
M1	J_{12}			J_4	J_1

$$L = (J_{10}, J_{11}, J_{12}, J_{13}, J_9, J_8, J_7, J_6, J_5, J_4, J_3, J_2, J_1)$$

Level 3
Level 2
Level 1

HLF/CP algorithm

- **Time complexity**

$O(|V|+|E|)$ ($|V|$ is the number of jobs and $|E|$ is the number of edges in the precedence graph)

- **Theorem (Hu, 1961) : HLF/CP for Tree**

- The HLF algorithm is optimal for $P_m \mid p_j = 1$, in-tree (out-tree) $\mid C_{\max}$.
- The HLF algorithm is optimal for $P_m \mid p_j = 1$, in-forest (out-forest) $\mid C_{\max}$.



HLF/CP algorithm

- N.F. Chen & C.L. Liu (1975)

The approximation ratio of HLF algorithm for the problem with general precedence constraints:

If $m = 2$, $\delta_{\text{HLF}} \leq 4/3$.

If $m \geq 3$, $\delta_{\text{HLF}} \leq 2 - 1/(m-1)$.

PTAS Algorithms: $P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$

- PTAS : Polynomial Time Approximation Scheme
- Approximation List scheduling policies
 - Graham's list algorithm/Greedy List
 - Discussed in Cilk Lectures : $T \leq 2T^*$, Also proved
 - CLR Book Chapter 27, Multi-threaded Algorithm
 - HLF algorithm
 - MSF algorithm

$$P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$$

Theorem 1

$P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$ is NP-complete.

1. Ullman (1976)

$$3\text{SAT} \leq P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$$

2. Lenstra and Rinnooy Kan (1978)

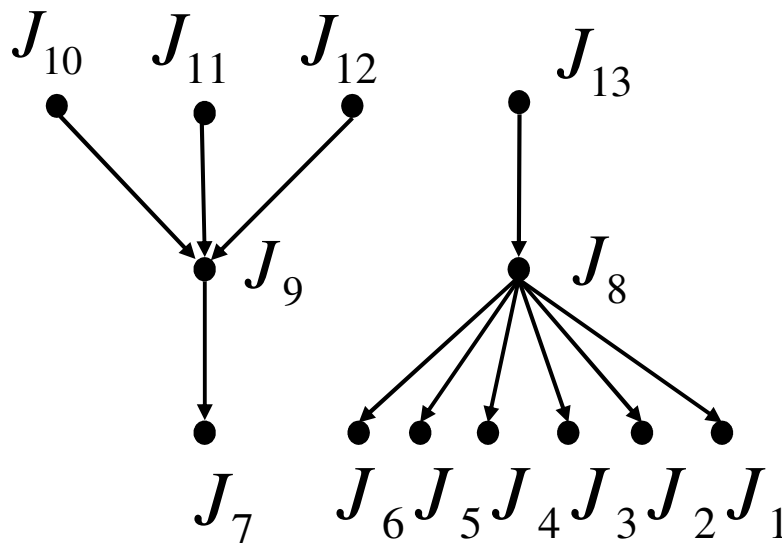
$$k\text{-clique} \leq P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$$

$P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$ is NP-complete.

Proof: out of Syllabus

List scheduling policies

- Set up a priority list L of jobs.
- When a processor is idle, assign the first ready job to the processor and remove it from the list L .



J_{11}	J_9	J_8	J_6	J_3
J_{10}	J_{13}	J_7	J_5	J_2
J_{12}			J_4	J_1

First job of the list
may not be ready

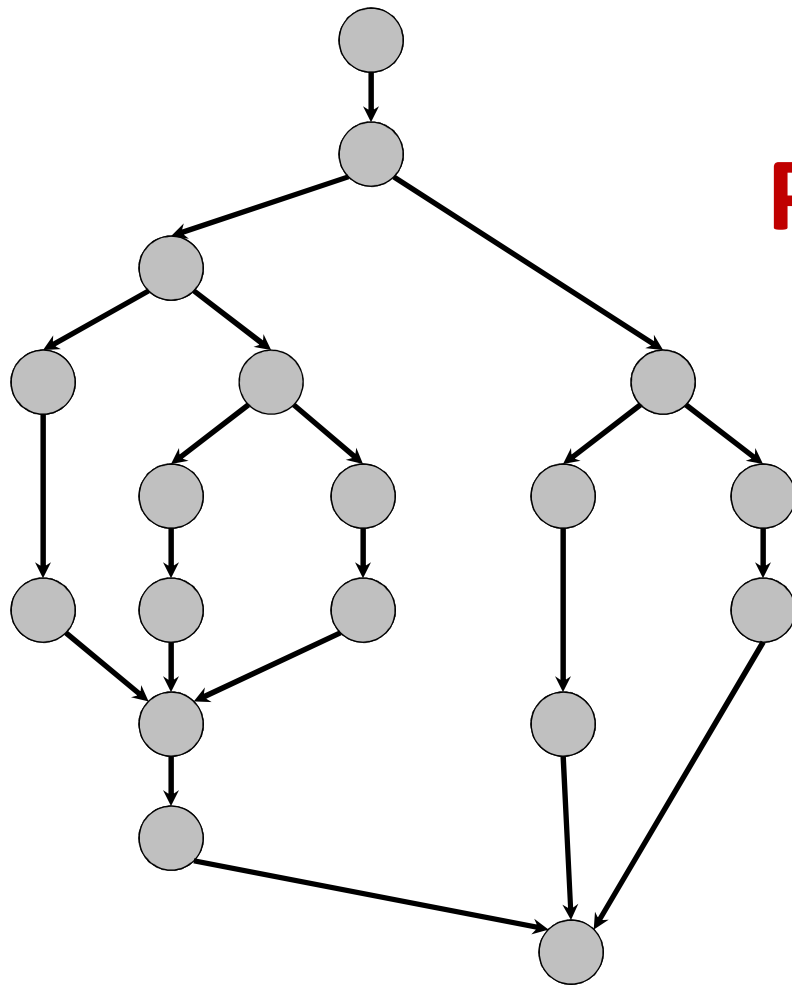
$$L = (J_9, J_8, J_7, J_6, J_5, J_{11}, J_{10}, J_{12}, J_{13}, J_4, J_3, J_2, J_1)$$

Graham's list algorithm

- Graham first analyzed the performance of the simplest list scheduling algorithm.
- List scheduling algorithm with an arbitrary job list is called Graham's list algorithm.
- Approximation ratio for $P_m \mid \text{prec}, p_j = 1 \mid C_{\max}$
 $\delta = 2 - 1/m$. (Tight bound!)
 - Approximation ratio is δ if for each input instance, the makespan produced by the algorithm is at most δ times of the optimal makespan.

CP Algo: CLR Book Page 779-783


T_P = execution time on P processors



$P_m \mid p_j = 1, \text{ prec} \mid C_{\max}$

CP Algorithms

T_P = execution time on P processors



$T_1 = \text{work}$

$T_\infty = \text{span}^*$

LOWER BOUNDS

- $T_P \geq T_1/P$
- $T_P \geq T_\infty$

* Also called *critical-path length* or *computational depth*.

CP: Greedy-Scheduling Theorem

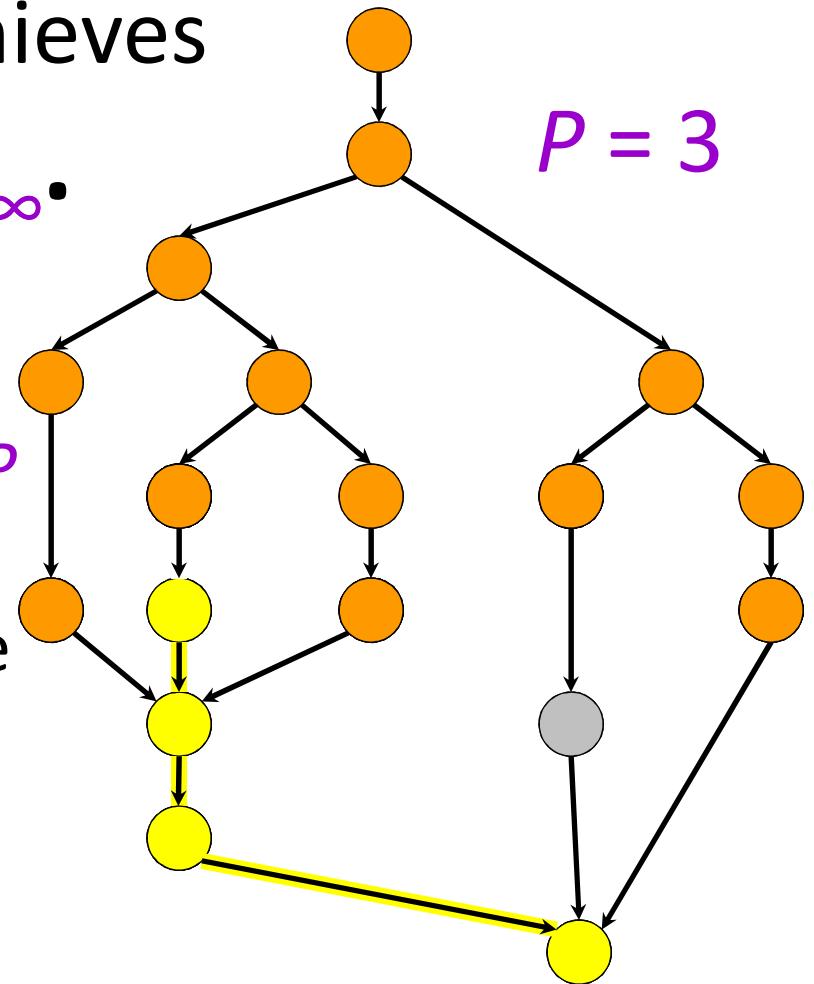
Theorem [Graham '68 & Brent '75].

Any greedy scheduler achieves

$$T_P \leq T_1/P + T_\infty.$$

Proof.

- # complete steps $\leq T_1/P$, since each complete step performs P work.
- # incomplete steps $\leq T_\infty$, since each incomplete step reduces the span of the unexecuted dag by 1. ■



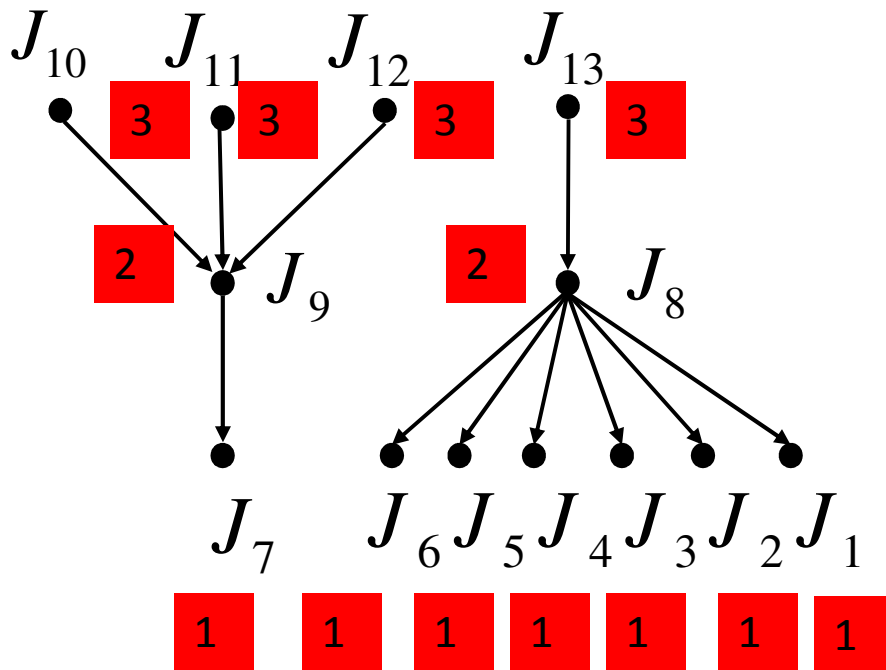
CP: Optimality of Greedy

Corollary. Any greedy scheduler achieves within a factor of 2 of optimal.

Proof. Let T_p^* be the execution time produced by the optimal scheduler. Since $T_p^* \geq \max\{T_1/P, T_\infty\}$ (lower bounds), we have

$$\begin{aligned} T_p &\leq T_1/P + T_\infty \\ &\leq 2 \cdot \max\{T_1/P, T_\infty\} \\ &\leq 2T_p^* . \quad \blacksquare \end{aligned}$$

HLF/CP algorithm : Example



M2	J_{10}	J_{13}	J_8	J_6	J_3
M2	J_{11}	J_9	J_7	J_5	J_2
M1	J_{12}			J_4	J_1

$$L = (J_{10}, J_{11}, J_{12}, J_{13}, J_9, J_8, J_7, J_6, J_5, J_4, J_3, J_2, J_1)$$

Level 3: $J_{10}, J_{11}, J_{12}, J_{13}$
 Level 2: J_9, J_8
 Level 1: $J_7, J_6, J_5, J_4, J_3, J_2, J_1$

HLF/CP algorithm

- **Time complexity**

$O(|V|+|E|)$ ($|V|$ is the number of jobs and $|E|$ is the number of edges in the precedence graph)

- **Theorem (Hu, 1961) : HLF/CP for Tree**

- The HLF algorithm is optimal for $P_m \mid p_j = 1$, in-tree (out-tree) $\mid C_{\max}$.
- The HLF algorithm is optimal for $P_m \mid p_j = 1$, in-forest (out-forest) $\mid C_{\max}$.



HLF/CP algorithm

- N.F. Chen & C.L. Liu (1975)

The approximation ratio of HLF algorithm for the problem with general precedence constraints:

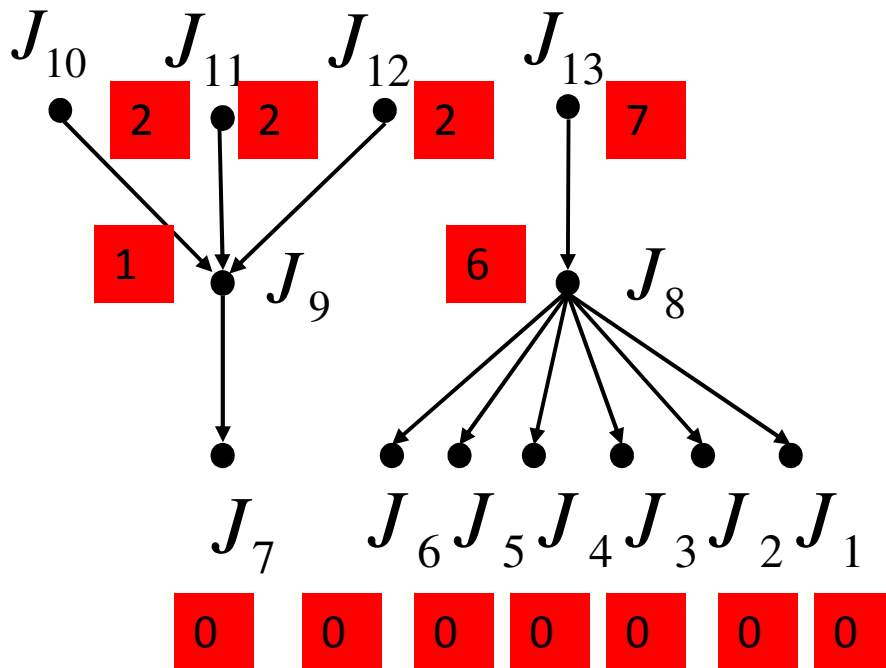
If $m = 2$, $\delta_{\text{HLF}} \leq 4/3$.

If $m \geq 3$, $\delta_{\text{HLF}} \leq 2 - 1/(m-1)$.

Most Successors First (MSF)

- Algorithm:
 - Set up a priority list L by nonincreasing order of the jobs' successors numbers.
 - (i.e. the job having more successors should have a higher priority in L than the job having fewer successors)
 - Execute the list scheduling policy based on this priority list L.

Most Successors First algorithm



M2	J_{13}	J_{10}	J_9	J_7	J_2
M2	J_{12}	J_8	J_6	J_4	J_1
M1	J_{11}		J_5	J_3	

$$L = (J_{13}, J_8, J_{12}, J_{11}, J_{10}, J_9, J_7, J_6, J_5, J_4, J_3, J_2, J_1)$$

7 6 2 2 2 1 0 0 0 0 0 0 0

Energy/Power/Temp Aware Scheduling of Tasks

Outline

- Power Aware
- Task with Hard Deadlines
- Energy Efficiency
- Energy Efficient Scheduling
- Real Time Tasks

Power Aware Scheduling Vs Energy Aware Scheduling

- Power Budget should not exceed
 - Minimized
 - Monthly Expenses: CAP ==> Solution is EMI
 - Power CAP: If your system have 100W design, at any instance of time you should not run things above 100W
 - Suppose you have 3KW wiring in your home, you have 3 AC with each of 1.5KW rating, At a given time, you can run maximum of 2 AC.
- Total energy budget should not exceed
 - Battery capacity, mah (mobile), AH (UPS)
 - Minimized: EC
 - Power and Time